

 北京大学计算机学院本科生课程

计算机组成与系统结构 实习




RISC-V ISA

北京大学微处理器研发中心

易江芳


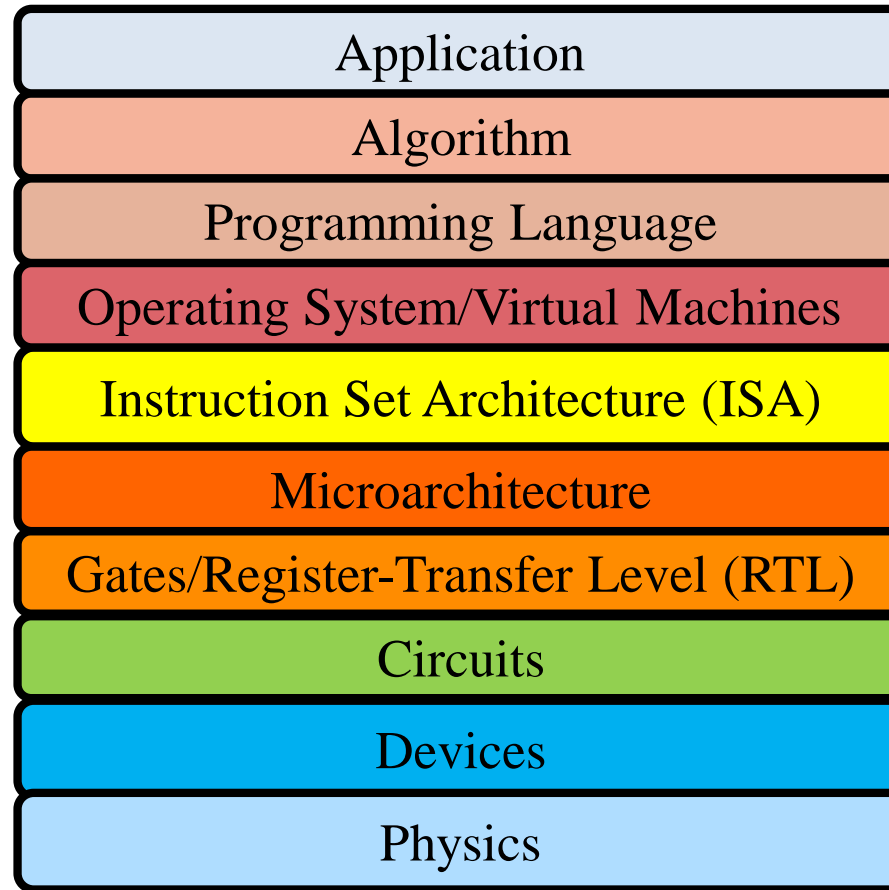
2022-10-10

Computer Architecture



In its broadest definition, computer architecture is the *design of the abstraction layers* that allow us to implement information processing applications efficiently using available manufacturing technologies.

ISA: software/hardware interface



ISA is not only the set of instructions

- ISA不仅仅是指令的集合，还包括了存储管理、异常中断处理、软硬件协同等多个方面
- 学习和设计ISA，需要建立高级语言与ISA相互映射的系统观
 - 算术语句 vs. 运算指令
 - 循环结构 vs. 分支转移
 - Switch结构 vs. 跳转
 - 静态全局变量 vs. 地址空间的静态数据区
 - 局部变量 vs. 栈
 - 动态分配存储 vs. 堆
 - 进程切换 vs. 存储管理和文件管理
 - 键盘和鼠标 vs. 中断处理

■ ■ ■ RISC ISA

- RISC philosophy
 - IBM 801
 - Cocke IBM, Patterson UCB, Hennessy Stanford, 1980s) – *Reduced Instruction Set Computing*
 - IBM PowerPC、 RISC I / RISC II、 MIPS etc.
 - Keep the instruction set small and simple, in order to build fast hardware
 - Let software do complicated operations by composing simpler ones

|| RISC-V ISA

- Fifth generation of RISC design from UC Berkeley
- Realistic & complete ISA, but open & simple
- Not over-architected for a certain implementation style
- Both 32-bit and 64-bit address space variants
 - RV32 and RV64
- Easy to subset/extend for education/research
 - RV32IM, RV32IMA, RV32IMAFD, RV32G
 - RV64IM, RV64IMA, RV64IMAFD, RV64G
- Techreport with RISC-V spec available on class website or riscv.org
- We' ll be using 64-bit RISC-V this semester in labs. Similar to MIPS you saw in Computer architecture.



RISC-V simple green card


RISC-V Reference Data				ARITHMETIC CORE INSTRUCTION SET			
RV64M Multiple Extension				RV64M Floating-Point Extensions			
MNEMONIC	NAME	DESCRIPTION (in Verilog)	NOTE	MNEMONIC	EXTNAME	DESCRIPTION (in Verilog)	NOTE
add rd, rs1, rs2	R ADD (Word)	R[rd] ← R[rs1] + R[rs2]	1)	mulh rd, rs1, rs2	R MULhpy (Word)	R[rd] ← (R[rs1] * R[rs2]) >> 32	1)
addi rd, rs1, imm	R ADDI (Word)	R[rd] ← R[rs1] + imm	1)	mulhu rd, rs1, rs2	R MULhpy (Upper Half)	R[rd] ← (R[rs1] * R[rs2]) >> 32	6)
addiw rd, rs1, imm	R ADDI (Word)	R[rd] ← R[rs1] + imm	1)	mulhs rd, rs1, rs2	R MULhpy (Lower Half)	R[rd] ← (R[rs1] * R[rs2]) >> 32	6)
and rd, rs1, rs2	R AND	R[rd] ← R[rs1] & R[rs2]	1)	div rd, rs1, rs2	R DIVd (Word)	R[rd] ← R[rs1] / R[rs2]	1)
andi rd, rs1, imm	R ANDI (Word)	R[rd] ← R[rs1] & imm	1)	divu rd, rs1, rs2	R DIVd (Upper Half)	R[rd] ← R[rs1] / R[rs2]	1)
auipc rd, rs1, imm	U Add Upper Immediate to PC	R[rd] ← PC + {imm, 12'b0}	1)	rem rd, rs1, rs2	R REMd (Word)	R[rd] ← R[rs1] % R[rs2]	1)
beq rs1, rs2, offset	SB Branch Equal	if (R[rs1] == R[rs2]) PC ← PC + {imm, 1b'0}	1)	remu rd, rs1, rs2	R REMd (Upper Half)	R[rd] ← R[rs1] % R[rs2]	1)
bge rs1, rs2, offset	SB Branch Greater than or Equal	if (R[rs1] >= R[rs2]) PC ← PC + {imm, 1b'0}	1)	sra rd, rs1, rs2	R SARd (Word)	R[rd] ← R[rs1] >> R[rs2]	1)
bgeu rs1, rs2, offset	SB Branch Greater than or Equal Unsigned	if (R[rs1] >= R[rs2]) PC ← PC + {imm, 1b'0}	2)	srai rd, rs1, rs2	R SARd (Upper Half)	R[rd] ← R[rs1] >> R[rs2]	1)
blt rs1, rs2, offset	SB Branch Less Than	if (R[rs1] < R[rs2]) PC ← PC + {imm, 1b'0}	2)	or rd, rs1, rs2	R ORd (Word)	R[rd] ← R[rs1] R[rs2]	1)
bltu rs1, rs2, offset	SB Branch Less Than Unsigned	if (R[rs1] < R[rs2]) PC ← PC + {imm, 1b'0}	2)	rem rd, rs1, rs2	R REMd (Word)	R[rd] ← R[rs1] % R[rs2]	1)
bne rs1, rs2, offset	SB Branch Not Equal	if (R[rs1] != R[rs2]) PC ← PC + {imm, 1b'0}	1)	and rd, rs1, rs2	R ANDd (Word)	R[rd] ← R[rs1] & R[rs2]	1)
csrrw rd, CSR, CSR	1 Cont./Stat.Reg.Read&Clear	R[rd] ← CSR; CSR ← CSR & ~imm	1)	lb rd, offset(rs1)	R LBd (Word)	R[rd] ← SignExt(Mem(R[rs1] + offset, byte))	1)
csrrs rd, CSR, CSR	1 Cont./Stat.Reg.Read&Set	R[rd] ← CSR; CSR ← CSR imm	1)	lh rd, offset(rs1)	R LHD (Word)	R[rd] ← SignExt(Mem(R[rs1] + offset, half))	1)
csrrwi rd, CSR, CSR	1 Cont./Stat.Reg.Read&Write	R[rd] ← CSR; CSR ← R[rs1]	1)	lw rd, offset(rs1)	R LWD (Word)	R[rd] ← Mem(R[rs1] + offset, word)	1)
csrrwi rd, CSR, CSR	1 Cont./Stat.Reg.Read&Write	R[rd] ← CSR; CSR ← imm	1)	ld rd, offset(rs1)	R LDD (Word)	R[rd] ← Mem(R[rs1] + offset, doubleword)	1)
ebreak	1 Environment Break	Transfer control to debugger	1)	addi rd, rs1, imm	R ADDI (Word)	R[rd] ← R[rs1] + imm	1)
ecall	1 Environment Call	Transfer control to operating system	1)	slli rd, rs1, imm	R SLLI (Word)	R[rd] ← R[rs1] << imm	1)
fence	1 Sync Thread	Synchronizes threads	1)	slti rd, rs1, imm	R SLTI (Word)	R[rd] ← (R[rs1] < imm) ? 1 : 0	1)
fence.i	1 Sync Inst. & Data	Synchronizes writes to instruction stream	1)	xori rd, rs1, imm	R XORI (Word)	R[rd] ← R[rs1] ^ imm	1)
jal	UJ Jump & Link	R[rd] ← PC + 4; PC ← PC + {imm, 1b'0}	1)	srl rd, rs1, imm	R SRLI (Word)	R[rd] ← R[rs1] >> imm	1)
jalr	1 Jump & Link Register	R[rd] ← PC + 4; PC ← R[rs1] + imm	3)	srai rd, rs1, imm	R SRAI (Word)	R[rd] ← R[rs1] >> imm	1)
lb	1 Load Byte	R[rd] ← {56'M[R[rs1] + imm] >> 7, 0}	4)	ori rd, rs1, imm	R ORI (Word)	R[rd] ← R[rs1] imm	1)
lbh	1 Load Byte Unsigned	R[rd] ← {56'M[R[rs1] + imm] >> 7, 0}	4)	andi rd, rs1, imm	R ANDI (Word)	R[rd] ← R[rs1] & imm	1)
ld	1 Load Doubleword	R[rd] ← {112'M[R[rs1] + imm] >> 15, 0}	4)	addiw rd, rs1, imm	R ADDIW (Word)	R[rd] ← R[rs1] + imm	1)
ldh	1 Load Halfword	R[rd] ← {48'M[R[rs1] + imm] >> 15, 0}	4)	jalr rd, rs1, imm	R JALR (Word)	R[rd] ← PC + 4	1)
ldu	1 Load Halfword Unsigned	R[rd] ← {48'M[R[rs1] + imm] >> 15, 0}	4)	ecall	R ECALL (Word)	PC ← R[rs1] + {imm, 1b'0}	1)
lui	U Load Upper Immediate	R[rd] ← {32'M[R[rs1] + imm] >> 15, 0}	4)	sb rs2, offset(rs1)	R SBD (Word)	Mem(R[rs1] + offset) ← R[rs2][7:0]	1)
luiw	U Load Word	R[rd] ← {32'M[R[rs1] + imm] >> 31, 0}	4)	sh rs2, offset(rs1)	R SHD (Word)	Mem(R[rs1] + offset) ← R[rs2][15:0]	1)
lw	1 Load Word Unsigned	R[rd] ← {32'M[R[rs1] + imm] >> 31, 0}	4)	sw rs2, offset(rs1)	R SWD (Word)	Mem(R[rs1] + offset) ← R[rs2][31:0]	1)
or	R OR	R[rd] ← R[rs1] R[rs2]	1)	sd rs2, offset(rs1)	R STD (Word)	Mem(R[rs1] + offset) ← R[rs2][63:0]	1)
ori	R ORI (Word)	R[rd] ← R[rs1] imm	1)	beq rs1, rs2, offset	R BEQ (Word)	if (R[rs1] == R[rs2]) PC ← PC + {offset, 1b'0}	1)
sb	S Store Byte	M[R[rs1] + imm] << 7 = R[rs2][7:0]	1)	bne rs1, rs2, offset	R BNE (Word)	if (R[rs1] != R[rs2]) PC ← PC + {offset, 1b'0}	1)
sh	S Store Halfword	M[R[rs1] + imm] << 15 = R[rs2][15:0]	1)	blt rs1, rs2, offset	R BLT (Word)	if (R[rs1] < R[rs2]) PC ← PC + {offset, 1b'0}	1)
shw	S Store Word	M[R[rs1] + imm] << 31 = R[rs2][31:0]	1)	bltu rs1, rs2, offset	R BLTU (Word)	if (R[rs1] < R[rs2]) PC ← PC + {offset, 1b'0}	1)
sl	R Set Less Than	R[rd] ← (R[rs1] < R[rs2]) ? 1 : 0	1)	bge rs1, rs2, offset	R BGE (Word)	if (R[rs1] >= R[rs2]) PC ← PC + {offset, 1b'0}	1)
slti	R Set Less Than Immediate	R[rd] ← (R[rs1] < imm) ? 1 : 0	1)	auipc rd, offset	R AUIPC (Word)	R[rd] ← PC + {offset, 12'b0}	1)
sltiu	R Set Less Than Immediate Unsigned	R[rd] ← (R[rs1] < imm) ? 1 : 0	1)	lui rd, offset	R LUI (Word)	R[rd] ← {offset, 12'b0}	1)
slu	R Set Less Than Unsigned	R[rd] ← (R[rs1] < R[rs2]) ? 1 : 0	2)	jal rd, imm	R JAL (Word)	R[rd] ← PC + 4	1)
sra	R Shift Right Arithmetic (Word)	R[rd] ← R[rs1] >> R[rs2]	1, 5)			PC ← PC + {imm, 1b'0}	1)
srai	R Shift Right Arithmetic (Upper Half)	R[rd] ← R[rs1] >> R[rs2]	1, 5)				
srl	R Shift Right (Word)	R[rd] ← R[rs1] >> R[rs2]	1)				
srlu	R Shift Right (Upper Half)	R[rd] ← R[rs1] >> R[rs2]	1)				
sub rd, rs1, rs2	R SUBd (Word)	R[rd] ← R[rs1] - R[rs2]	1)				
subw rd, rs1, rs2	R SUBd (Word)	R[rd] ← R[rs1] - R[rs2]	1)				
sw	S Store Word	M[R[rs1] + imm] << 31 = R[rs2][31:0]	1)				
xor	R XOR	R[rd] ← R[rs1] ^ R[rs2]	1)				
xori	R XORI (Word)	R[rd] ← R[rs1] ^ imm	1)				

RISC-V RV64I Simple Green Card

Instruction	Type	Opcode	Func3	Func7/IMM	Operation
add rd, rs1, rs2	R	0x33	0x0	0x00	R[rd] ← R[rs1] + R[rs2]
mul rd, rs1, rs2			0x0	0x01	R[rd] ← (R[rs1] * R[rs2])[31:0]
sub rd, rs1, rs2			0x0	0x20	R[rd] ← R[rs1] - R[rs2]
sll rd, rs1, rs2			0x1	0x00	R[rd] ← R[rs1] << R[rs2]
mulh rd, rs1, rs2			0x1	0x01	R[rd] ← (R[rs1] * R[rs2])[63:32]
slt rd, rs1, rs2			0x2	0x00	R[rd] ← (R[rs1] < R[rs2]) ? 1 : 0
xor rd, rs1, rs2			0x4	0x00	R[rd] ← R[rs1] ^ R[rs2]
div rd, rs1, rs2			0x4	0x01	R[rd] ← R[rs1] / R[rs2]
srl rd, rs1, rs2	I	0x03	0x5	0x00	R[rd] ← R[rs1] >> R[rs2]
sra rd, rs1, rs2			0x5	0x20	R[rd] ← R[rs1] >> R[rs2]
or rd, rs1, rs2			0x6	0x00	R[rd] ← R[rs1] R[rs2]
rem rd, rs1, rs2			0x6	0x01	R[rd] ← (R[rs1] % R[rs2])
and rd, rs1, rs2			0x7	0x00	R[rd] ← R[rs1] & R[rs2]
lb rd, offset(rs1)			0x0		R[rd] ← SignExt(Mem(R[rs1] + offset, byte))
lh rd, offset(rs1)			0x1		R[rd] ← SignExt(Mem(R[rs1] + offset, half))
lw rd, offset(rs1)			0x2		R[rd] ← Mem(R[rs1] + offset, word)
ld rd, offset(rs1)			0x3		R[rd] ← Mem(R[rs1] + offset, doubleword)
addi rd, rs1, imm	I	0x13	0x0		R[rd] ← R[rs1] + imm
slli rd, rs1, imm			0x1	0x00	R[rd] ← R[rs1] << imm
slti rd, rs1, imm			0x2		R[rd] ← (R[rs1] < imm) ? 1 : 0
xori rd, rs1, imm			0x4		R[rd] ← R[rs1] ^ imm
srl rd, rs1, imm			0x5	0x00	R[rd] ← R[rs1] >> imm
srai rd, rs1, imm			0x5	0x20	R[rd] ← R[rs1] >> imm
ori rd, rs1, imm			0x6		R[rd] ← R[rs1] imm
andi rd, rs1, imm			0x7		R[rd] ← R[rs1] & imm
addiw rd, rs1, imm	I	0x1B	0x0		R[rd] ← R[rs1] + imm
jalr rd, rs1, imm			0x67	0x0	PC ← R[rs1] + {imm, 1b'0}
ecall			0x0	0x000	(Transfers control to operating system)
					a0 = 1 is print value of a1 as an integer.
					a0 = 10 is exit or end of code indicator.
sb rs2, offset(rs1)	S	0x23	0x0		Mem(R[rs1] + offset) ← R[rs2][7:0]
sh rs2, offset(rs1)			0x1		Mem(R[rs1] + offset) ← R[rs2][15:0]
sw rs2, offset(rs1)			0x2		Mem(R[rs1] + offset) ← R[rs2][31:0]
sd rs2, offset(rs1)			0x3		Mem(R[rs1] + offset) ← R[rs2][63:0]
beq rs1, rs2, offset			0x0		if (R[rs1] == R[rs2]) PC ← PC + {offset, 1b'0}
bne rs1, rs2, offset			0x1		if (R[rs1] != R[rs2]) PC ← PC + {offset, 1b'0}
blt rs1, rs2, offset			0x4		if (R[rs1] < R[rs2]) PC ← PC + {offset, 1b'0}
bltu rs1, rs2, offset			0x5		if (R[rs1] < R[rs2]) PC ← PC + {offset, 1b'0}
bge rs1, rs2, offset	U	0x17	0x0		if (R[rs1] >= R[rs2]) PC ← PC + {offset, 1b'0}
auipc rd, offset			0x17		R[rd] ← PC + {offset, 12'b0}
lui rd, offset			0x37		R[rd] ← {offset, 12'b0}
jal rd, imm			0x6f		R[rd] ← PC + 4
					PC ← PC + {imm, 1b'0}

RISC-V Instruction Assembly Syntax

- Instructions have an opcode and operands
- E.g., `add x1, x2, x3 # x1 = x2 + x3`



•Assembly

– Example: `add x1, x2, x3` (in RISC-V)

– Equivalent to: `a = b + c` (in C)

where C variables \Leftrightarrow RISC-V registers are : $a \Leftrightarrow x1$, $b \Leftrightarrow x2$, $c \Leftrightarrow x3$

Summary of RISC-V Instruction Formats

Additional opcode bits/immediate				Source Reg. 2		Source Reg. 1		Destination Reg.		7-bit opcode field (but low 2 bits =11 ₂)					
31	30	25	24	21	20	19	15	14	12	11	8	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type	
imm[11:0]						rs1		funct3		rd		opcode		I-type	
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type	
imm[12]	imm[10:5]		rs2		rs1		funct3		imm[4:1]	imm[11]	opcode		B-type		
imm[31:12]									rd		opcode		U-type		
imm[20]	imm[10:1]			imm[11]	imm[19:12]				rd		opcode		J-type		

- Aligned on a four-byte boundary in memory
- Sign bit of immediates always on bit 31 of instruction.
- Register fields never move
- Opcode fields is on the right

Where is NOP?

- `Addi x0, x0, #0`


RISC-V Instruction Encoding

xxxxxxxxxxxxxxxxaa			16-bit ($aa \neq 11$)
xxxxxxxxxxxxxxxx	xxxxxxxxxxxxbbb11	32-bit ($bbb \neq 111$)	
...XXXX	xxxxxxxxxxxxxxxx	xxxxxxxxxxx011111	48-bit
...XXXX	xxxxxxxxxxxxxxxx	xxxxxxxxxxx011111	64-bit
...XXXX	xxxxxxxxxxxxxxxx	xxxxnnnn111111	$(80+16*nnnn)$ -bit, $nnnn \neq 1111$
...XXXX	xxxxxxxxxxxxxxxx	xxxxx1111111111	Reserved for ≥ 320 -bits

- Base instruction set (RV32 and RV64) always has fixed 32-bit instructions lowest two bits = 11_2
- All branches and jumps have targets at 16-bit granularity (even in base ISA where all instructions are fixed 32 bits)
 - Still will cause a fault if fetching a 32-bit instruction

■ ■ ■ Data formats and addressing

- Data formats
 - 8b Bytes, 16b Half words, 32b words and 64b double words
- Some issues
 - Byte addressing
 - Little endian
 - Word alignment



Example: RISC-V rv32I/64I R-format instructions

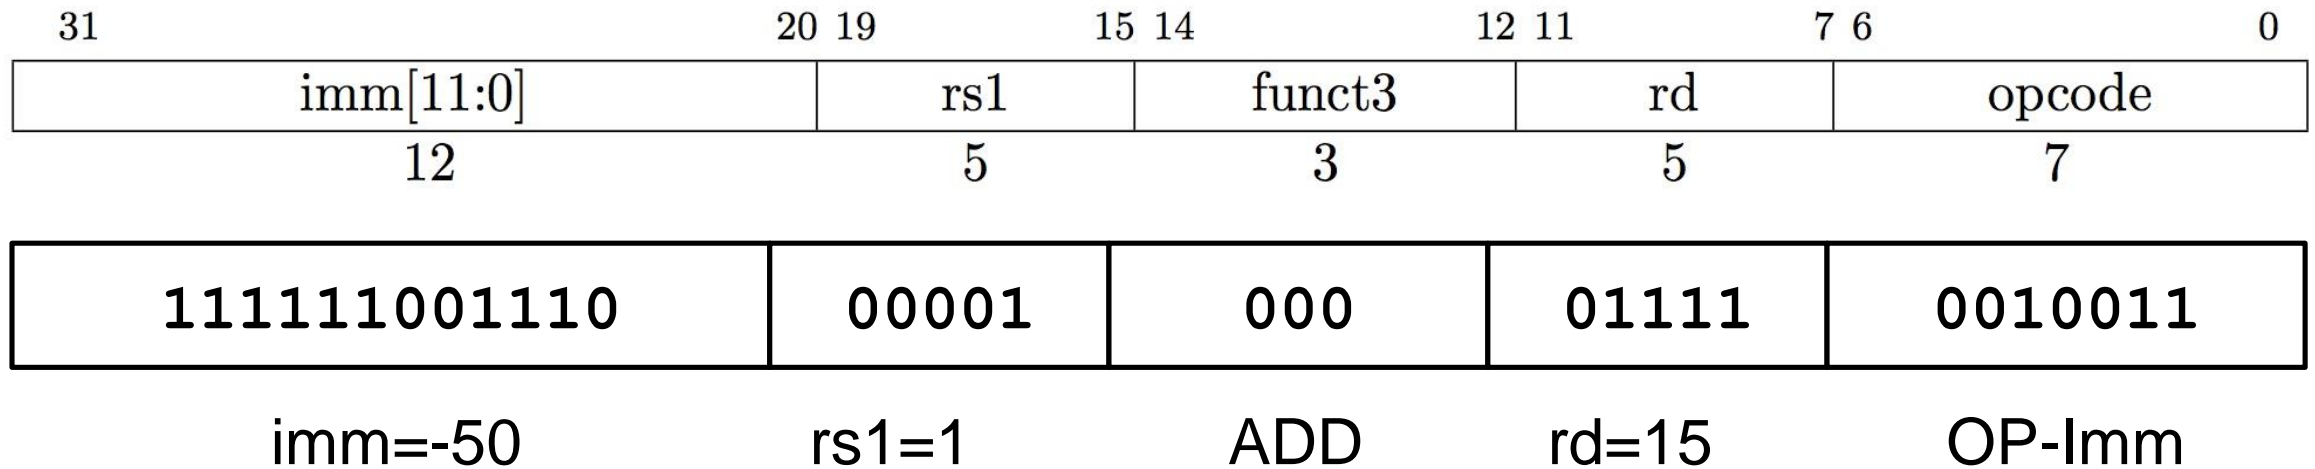
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

Different encoding in funct7 + funct3 selects different operations

RISC-V I-Format Instruction

- RISC-V Assembly Instruction:


addi x15, x1, -50



- imm[11:0] can hold values in range $[-2048_{10}, +2047_{10}]$
- Immediate is always sign-extended to 32-bits before use in an arithmetic operation

Support 32-bit ops in rv64I


- Integer Register-Immediate Instructions
 - ADDIW
 - $Rd = (\text{sign extend})(rs1 + \text{imm12})$
 - Ignore overflow




– Shift instruction

31	26	25	24	20	19	15	14	12	11	7	6	0
imm[11:6]			imm[5]	imm[4:0]		rs1	funct3	rd		opcode		
6			1	5		5	3	5		7		
000000			shamt[5]	shamt[4:0]		src	SLLI	dest		OP-IMM		
000000			shamt[5]	shamt[4:0]		src	SRLI	dest		OP-IMM		
010000			shamt[5]	shamt[4:0]		src	SRAI	dest		OP-IMM		
000000			0	shamt[4:0]		src	SLLIW	dest		OP-IMM-32		
000000			0	shamt[4:0]		src	SRLIW	dest		OP-IMM-32		
010000			0	shamt[4:0]		src	SRAIW	dest		OP-IMM-32		

RISC-V B-Format for Branches



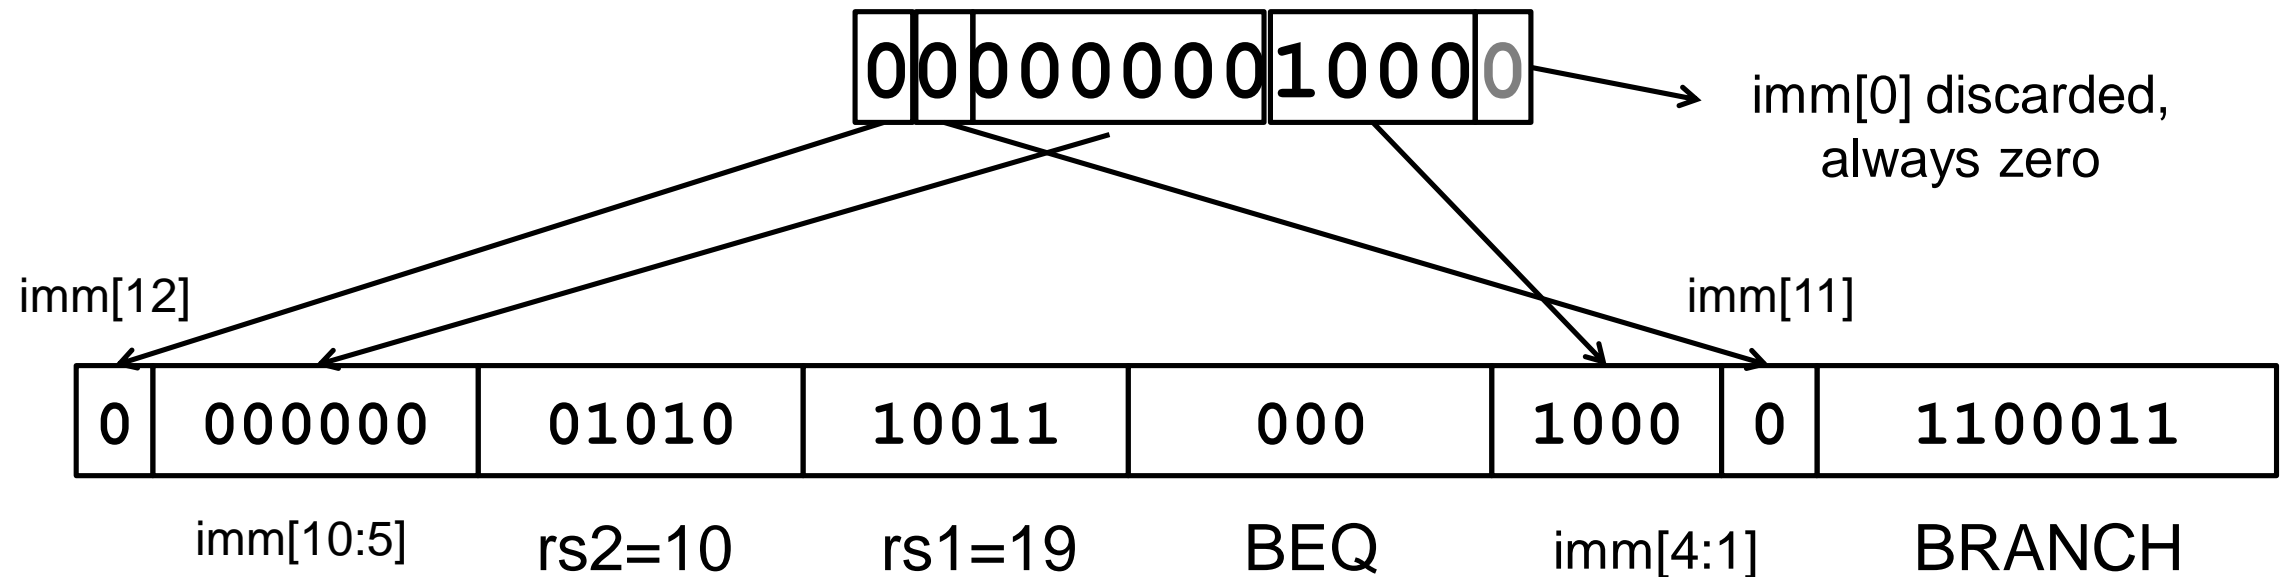
- B-format is mostly same as S-Format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)




Branch Example

`beq x19,x10, offset = 16 bytes`

13-bit immediate, imm[12:0], with value 16



U-Format for “Upper Immediate” instructions



- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - LUI – Load Upper Immediate
 - AUIPC – Add Upper Immediate to PC

|| LUI to create long immediates


- LUI writes the upper 20 bits of the destination with the immediate value, and clears the lower 12 bits.
- Together with an ADDI to set low 12 bits, can create any 32-bit value in a register using two instructions (LUI/ADDI).

LUI x10, 0x87654 # x10 = 0x87654000

ADDI x10, x10, 0x321 # x10 = 0x87654321

思考：如果要得到0xDEADBEEF，如何？

RISC-V J-Format for Jump Instructions

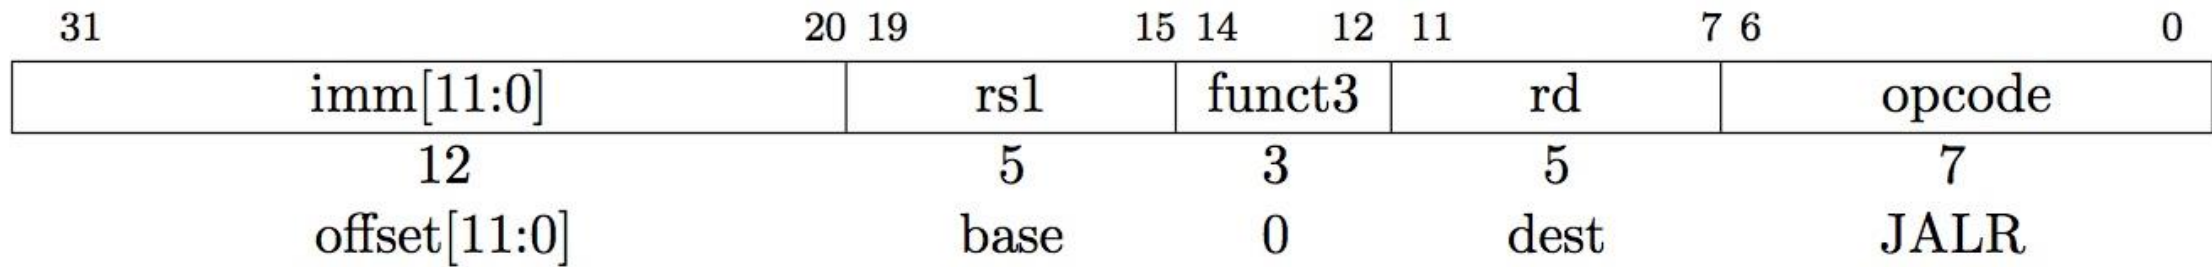


- JAL saves PC+4 in register rd (the return address)
 - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Example:

```
# j pseudo-instruction
j Label = jal x0, Label # Discard return address
```

```
# Call function within  $2^{18}$  instructions of PC
jal ra, FuncName
```

JALR Instruction (I-Format)



- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes

Example:

```
# ret and jr psuedo-instructions
ret = jr ra = jalr x0, ra, 0
# Call function at any 32-bit absolute address
lui x1, <hi20bits>
jalr ra, x1, <lo12bits>
# Jump PC-relative with 32-bit offset
auipc x1, <hi20bits>
jalr x0, x1, <lo12bits>
```

RISC-V Pseudo-instructions

PSEUDO INSTRUCTIONS

③

MNEMONIC	NAME	DESCRIPTION	USES
beqz	Branch = zero	$\text{if}(R[rs1]==0) PC=PC+\{imm,1b'0\}$	beq
bnez	Branch \neq zero	$\text{if}(R[rs1]!=0) PC=PC+\{imm,1b'0\}$	bne
fabs.s, fabs.d	Absolute Value	$F[rd] = (F[rs1] < 0) ? -F[rs1] : F[rs1]$	fsgnx
fmv.s, fmv.d	FP Move	$F[rd] = F[rs1]$	fsgnj
fneg.s, fneg.d	FP negate	$F[rd] = -F[rs1]$	fsgnjn
j	Jump	$PC = \{imm,1b'0\}$	jal
jr	Jump register	$PC = R[rs1]$	jalr
la	Load address	$R[rd] = \text{address}$	auipc
li	Load imm	$R[rd] = \text{imm}$	addi
mv	Move	$R[rd] = R[rs1]$	addi
neg	Negate	$R[rd] = -R[rs1]$	sub
nop	No operation	$R[0] = R[0]$	addi
not	Not	$R[rd] = \sim R[rs1]$	xori
ret	Return	$PC = R[1]$	jalr
seqz	Set = zero	$R[rd] = (R[rs1]==0) ? 1 : 0$	sltiu
snez	Set \neq zero	$R[rd] = (R[rs1]!=0) ? 1 : 0$	sltu

- 使用编译选项，可在编译用户程序时不生成伪指令：

– `riscv64-unknown-elf-objdump -Mno-aliases -D xxx`

|| RISC-V 标准扩展集

- M集：标准整数乘除扩展集，在整数集的基础上增加了乘法和除法运算
- A集：标准原子指令扩展集，使得两个处理器之间同步原子地读，修改和写memory
- F集：标准单精度浮点扩展集，有浮点的运算指令和load/store指令
- D集：标准双精度浮点扩展集，有双精度的运算指令和load/store指令
- Q集：标准四倍精度浮点扩展集，有四倍精度的运算指令和load/store指令
- L集：标准十进制浮点扩展集，用于64位，128位的浮点运算
- C集：标准压缩指令扩展集，将32位的指令能与16位的指令相混合
- B集：标准位操作扩展集，能插入，删除，测试比特域
- T集：标准内存交易扩展集，用于内存交易
- P集：标准Packed-SIMD指令集，用于打包的SIMD指令

Lab2.1中涉及的指令

RISC-V RV64I Simple Green Card

Instruction	Type	Opcode	Funct3	Funct7/IMM	Operation
add rd, rs1, rs2	R	0x33	0x0	0x00	$R[rd] \leftarrow R[rs1] + R[rs2]$
mul rd, rs1, rs2			0x0	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[31:0]$
sub rd, rs1, rs2			0x0	0x02	$R[rd] \leftarrow R[rs1] - R[rs2]$
sll rd, rs1, rs2			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll R[rs2]$
mulh rd, rs1, rs2			0x1	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
slt rd, rs1, rs2			0x2	0x00	$R[rd] \leftarrow (R[rs1] < R[rs2]) ? 1 : 0$
xor rd, rs1, rs2			0x4	0x00	$R[rd] \leftarrow R[rs1] \wedge R[rs2]$
div rd, rs1, rs2			0x4	0x01	$R[rd] \leftarrow R[rs1] / R[rs2]$
srl rd, rs1, rs2			0x5	0x00	$R[rd] \leftarrow R[rs1] \gg R[rs2]$
sra rd, rs1, rs2			0x5	0x20	$R[rd] \leftarrow R[rs1] \ggg R[rs2]$
or rd, rs1, rs2			0x6	0x00	$R[rd] \leftarrow R[rs1] \vee R[rs2]$
rem rd, rs1, rs2			0x6	0x01	$R[rd] \leftarrow (R[rs1] \% R[rs2])$
and rd, rs1, rs2			0x7	0x00	$R[rd] \leftarrow R[rs1] \& R[rs2]$
lb rd, offset(rs1)			0x0		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset}, \text{byte}))$
lh rd, offset(rs1)			0x1		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset}, \text{half}))$
lw rd, offset(rs1)			0x2		$R[rd] \leftarrow \text{Mem}(R[rs1] + \text{offset}, \text{word})$
ld rd, offset(rs1)			0x3		$R[rd] \leftarrow \text{Mem}(R[rs1] + \text{offset}, \text{doubleword})$
addi rd, rs1, imm	I	0x03	0x0		$R[rd] \leftarrow R[rs1] + \text{imm}$
slli rd, rs1, imm			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll \text{imm}$
slti rd, rs1, imm			0x2		$R[rd] \leftarrow (R[rs1] < \text{imm}) ? 1 : 0$
xori rd, rs1, imm			0x4		$R[rd] \leftarrow R[rs1] \wedge \text{imm}$
srl rd, rs1, imm		0x13	0x5	0x00	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
srai rd, rs1, imm			0x5	0x20	$R[rd] \leftarrow R[rs1] \ggg \text{imm}$
ori rd, rs1, imm			0x6		$R[rd] \leftarrow R[rs1] \vee \text{imm}$
andi rd, rs1, imm			0x7		$R[rd] \leftarrow R[rs1] \& \text{imm}$
addiw rd, rs1, imm		0x1B	0x0		$R[rd] \leftarrow \text{SignExt}(R[rs1][31:0] + \text{imm})$
jalr rd, rs1, imm		0x67	0x0		$R[rd] \leftarrow PC + 4$ $PC \leftarrow R[rs1] + \{\text{imm}, 1b'0\}$
ecall		0x73	0x0	0x000	(Transfers control to operating system)
					a0 = 1 is print value of a1 as an integer.
					a0 = 10 is exit or end of code indicator.
sb rs2, offset(rs1)	S	0x23	0x0		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][7:0]$
sh rs2, offset(rs1)			0x1		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][15:0]$
sw rs2, offset(rs1)			0x2		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][31:0]$
sd rs2, offset(rs1)			0x3		$\text{Mem}(R[rs1] + \text{offset}) \leftarrow R[rs2][63:0]$
beq rs1, rs2, offset	SB	0x63	0x0		if($R[rs1] == R[rs2]$) $PC \leftarrow PC + \{\text{offset}, 1b'0\}$
			0x1		if($R[rs1] != R[rs2]$) $PC \leftarrow PC + \{\text{offset}, 1b'0\}$
bne rs1, rs2, offset					
blt rs1, rs2, offset			0x4		if($R[rs1] < R[rs2]$) $PC \leftarrow PC + \{\text{offset}, 1b'0\}$
bge rs1, rs2, offset			0x5		if($R[rs1] \geq R[rs2]$) $PC \leftarrow PC + \{\text{offset}, 1b'0\}$
auipc rd, offset	U	0x17			$R[rd] \leftarrow PC + \{\text{offset}, 12'b0\}$
lui rd, offset		0x37			$R[rd] \leftarrow \{\text{offset}, 12'b0\}$
jal rd, imm	UJ	0x6f			$R[rd] \leftarrow PC + 4$ $PC \leftarrow PC + \{\text{imm}, 1b'0\}$

For further reference, here are the bit lengths of the instruction components

R-TYPE	funct7	rs2	rs1	funct3	rd	opcode
Bits	7	5	5	3	5	7

I-TYPE	imm[11:0]	rs1	funct3	rd	opcode
Bits	12	5	3	5	7

S-TYPE	imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode
Bits	7	5	5	3	5	7

SB-TYPE	imm[12]	imm[10:5]	rs2	rs1	funct3	imm[4:1]	imm[11]	opcode
Bits	1	6	5	5	3	4	1	7

U-TYPE	imm[31:12]	rd	opcode
Bits	20	5	7

UJ-TYPE	imm[20]	imm[10:1]	imm[11]	imm[19:12]	rd	opcode
Bits	1	10	1	8	5	7

执行结果参考:

<https://kvakil.github.io/venus/>

■ ■ ■ RISC-V Functional Call

1. Put parameters in a place where function can access them
2. Transfer control to function
3. Acquire (local) storage resources needed for function
4. Perform desired task of the function
5. Put result value in a place where calling code can access it and restore any registers you used
6. Return control to point of origin, since a function can be called from several points in a program

■ ■ ■ RISC-V calling convention register usage

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5–7	t0–2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	s1	Saved register	Callee
x10–11	a0–1	Function arguments/return values	Caller
x12–17	a2–7	Function arguments	Caller
x18–27	s2–11	Saved registers	Callee
x28–31	t3–6	Temporaries	Caller
f0–7	ft0–7	FP temporaries	Caller
f8–9	fs0–1	FP saved registers	Callee
f10–11	fa0–1	FP arguments/return values	Caller
f12–17	fa2–7	FP arguments	Caller
f18–27	fs2–11	FP saved registers	Callee
f28–31	ft8–11	FP temporaries	Caller

■ ■ ■ Optimized Function Convention

To reduce expensive loads and stores from spilling and restoring registers, RISC-V function-calling convention divides registers into two categories:


1. Preserved across function call
 - Caller can rely on values being unchanged
 - `sp, gp, tp`, “saved registers” `s0- s11` (`s0` is also `fp`)
2. Not preserved across function call
 - Caller *cannot* rely on values being unchanged
 - Argument/return registers `a0-a7, ra`, “temporary registers” `t0- t6`

|| Allocating Space on Stack

- C has two storage classes: automatic and static
 - *Automatic* variables are local to function and discarded when function exits
 - *Static* variables exist across exits from and entries to procedures
- Use stack for automatic (local) variables that don't fit in registers
- *Procedure frame* or *activation record*: segment of stack with saved registers and local variables

■ ■ ■ RISC-V RV32 memory map

- RV64 convention (RV32 and RV128 have different memory layouts)
- Stack starts in high memory and grows down
 - Hexadecimal (base 16) : `0xffff_ffff_hex`
 - Stack must be aligned on 16-byte boundary (not true in examples above)
- RV64 programs (*text segment*) in low end
 - `0x0001_0000_hex`
- *static data segment* (constants and other static variables) above text for static variables
 - RISC-V convention *global pointer* (`gp`) points to static
 - RV64 `gp` = `0x1000_0000_hex`
- *Heap* above static for data structures that grow and shrink ; grows up to high addresses



欢迎提问

系统调用的具体实现

- 内核中为每个系统调用定义了唯一的编号，就是系统调用号
- 内核中保存了一张系统调用表，该表中包含了系统调用号和其对应的服务例程地址。第n个表项包含系统调用号为n的服务例程地址。
- 系统调用陷入内核前，需要把系统调用号一起传入内核，而该号实际上是系统调用表的下标。
 - 在RISC-V模拟器中，这个传递动作可以通过在执行 `scall` 前把系统调用号装入寄存器 **a[7]** 实现，这样系统调用处理程序一旦运行，就可以从 `a[7]` 中得到系统调用号，然后再去系统调用表中中寻找相应服务例程
 - 通过什么方式传递系统调用号，由 **ABI** 决定

■ ■ ■ 系统调用的返回

- 以x86为例
- `system_call()`从**`eax`**获得系统调用的返回值, 并把这个值存放在曾保存用户态**`eax`**寄存器栈单元的那个位置上, 然后跳转到 `ret_from_sys_call()`, 终止系统调用处理程序的执行

■ ■ ■ Acknowledgements

- These slides contain material developed and copyright by:
 - Arvind (MIT)
 - Krste Asanovic (MIT/UCB)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)
- MIT material derived from course 6.823
- UCB material derived from course CS252

■ || RISC-V Function Call Conventions

- **a0–a7** (**x10–x17**): eight *argument* registers to pass parameters and two return values (**a0–a1**)
- **ra**: one *return address* register to return to the point of origin (**x1**)
- As a special function call, system calls also use **a0–a7** to transfer arguments.

不同指令集	系统调用指令	系统调用号 保存寄存器	系统调用参数保 存寄存器
X86(linux)	int \$0x80	eax	ebx, ecx, edx, esi, edi
RISCV(linux)	scall	a[7]	a[0]~a[3]

```
#define __NR_read 3  
#define __NR_write 4
```

系统调用号举例（x86）

```
#define SYS_read 63  
#define SYS_write 64
```

系统调用号举例（RISCV）

syscall_riscv.h

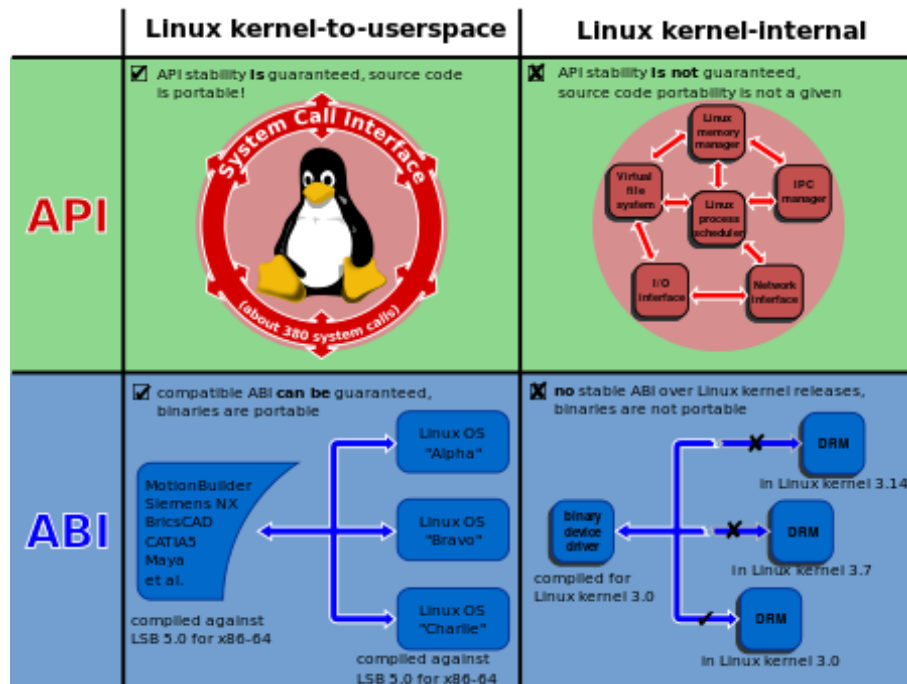
```
static inline long
__internal_syscall(long n, long _a0, long _a1, long _a2, long _a3)
{
    register long a0 asm("a0") = _a0;
    register long a1 asm("a1") = _a1;
    register long a2 asm("a2") = _a2;
    register long a3 asm("a3") = _a3;
    register long a7 asm("a7") = n;

    asm volatile ("syscall\n"
                  "bltz a0, __syscall_error"
                  : "+r"(a0) : "r"(a1), "r"(a2), "r"(a3), "r"(a7));

    return a0;
}
```


系统调用 (system call)

- 功能概述
 - 位于用户空间的应用程序与位于内核空间的操作系统内核程序之间的功能接口
- 作用
 - 使用户程序与内核程序相分离，保护内核，提高系统的安全性
 - 为用户提供有关设备管理、输入输出系统、文件系统和进程控制、通信及存储管理等方面的功能
 - 对用户隐藏系统程序的内部结构
 - 编程容易，从硬件设备的编程中解脱出来



系统调用是用户态进入内核态的唯一入口

一般的系统调用处理过程



系统调用处理过程^[1]

