

 北京大学计算机学院本科生课程

计算机组成与 系统结构实习



数据级并行

北京大学微处理器研发中心

2023-11-20

Outline

- DLP
- SIMD
- GPU

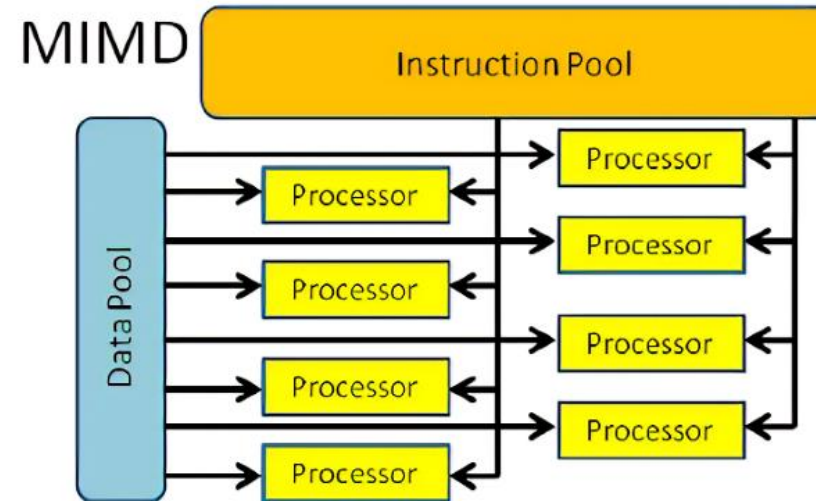

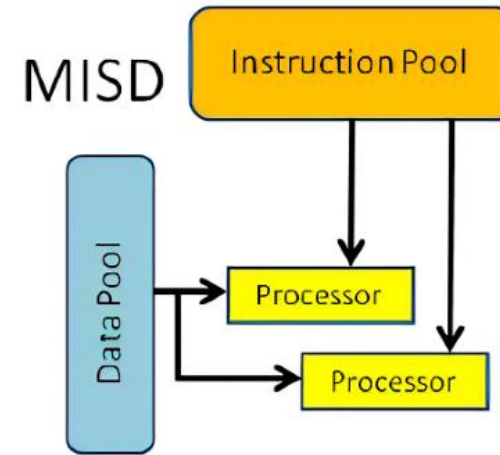

|| Flynn's Taxonomy

- Classified by data and control streams in 1966

Single Instruction Single Data (SISD) (Uniprocessor)	Single Instruction Multiple Data (SIMD) (single PC: Vector, CM-2, SSE)
Multiple Instruction Single Data (MISD) (????)	Multiple Instruction Multiple Data (MIMD) (Clusters, SMP servers)

- SISD \Rightarrow Implicit Instruction Level Parallelism (ILP)
- SIMD \Rightarrow Data Level Parallelism (DLP)
- MIMD \Rightarrow Thread Level Parallelism (TLP)

Flynn's Taxonomy



■ ■ ■ Data Level Parallelism (DLP)

We call these algorithms **data parallel** because their parallelism comes from simultaneous operations across large sets of data, rather than from multiple threads of control.

- W. Daniel Hillis and Guy L. Steele
"Data Level Algorithms", C. ACM

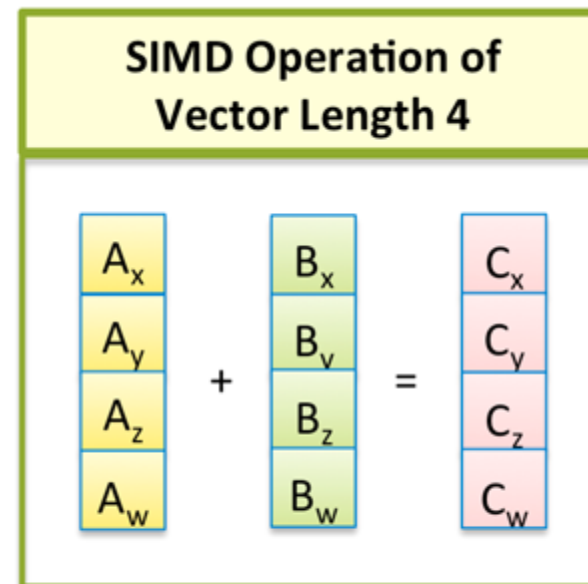

|| Vectors Motivation

```
for (int i = 0; i < N; i++)  
    A[i] = B[i] + C[i];
```

|| CPU Data-level Parallelism

- Let's make the execution unit (ALU) really wide
- Let's make the registers really wide too

```
for (int i = 0; i < N; i += 4) {  
    // in parallel  
    A[i] = B[i] + C[i];  
    A[i+1] = B[i+1] + C[i+1];  
    A[i+2] = B[i+2] + C[i+2];  
    A[i+3] = B[i+3] + C[i+3];  
}
```



Intel® Architecture currently has SIMD operations of vector length 4, 8, 16

Examples

- PyTorch

```
1 import torch
2
3 SIZE = 100
4 a = torch.rand(SIZE)
5 b = torch.rand(SIZE)
6 c = a+b
7 d = a.sum()
```

- Scala

```
val vector = Vector(1, 2, 3, 4, 5, 6)
val plusOne = vector.map(x => x + 1)
val sum = vector.reduce((x, y) => x + y)
```

- Functional programming
 - map , reduce
- Compared with PyTorch
 - $c=a+b$ -> 'map' (element-wise op)
 - $d=a.sum()$ -> 'reduce' (reducing a vector into a scalar)
- Scales to distributed systems

Examples

- X86 AVX intrinsics
 - Platform-dependent
 - Add (map operation)
 - `_m256` contains `float *8`
 - Sum (reduce operation)
 - Ignoring non-associativity of float

```
1  #include <immintrin.h>
2
3  constexpr int WIDTH = 256 / 8 / sizeof(float);
4
5  // c = a + b
6  // assumption:
7  //     1. a, b, c are aligned by 32 Bytes
8  //     2. N % WIDTH == 0
9  void vectorAdd(const float a[], const float b[], float c[], int N) {
10     for (auto i: int = 0; i < N; i += WIDTH) {
11         auto va: __m256 = _mm256_load_ps(p: &a[i]);
12         auto vb: __m256 = _mm256_load_ps(p: &b[i]);
13         auto vc: __m256 = va + vb;
14         _mm256_store_ps(p: &c[i], a: vc);
15     }
16 }
```

```
18 // assumption:
19 //     1. a is aligned by 32 Bytes
20 //     2. N % WIDTH == 0
21 float vectorSum(const float a[], int N) {
22     __m256 partialSum {};
23     for (auto i: int = 0; i < N; i += WIDTH) {
24         auto va: __m256 = _mm256_load_ps(p: &a[i]);
25         partialSum += va;
26     }
27     auto res: float = 0.0f;
28     // will be optimized to log(WIDTH) time by -Ofast
29     for (auto i: int = 0; i < WIDTH; i++) {
30         res += partialSum[i];
31     }
32 }
```

DLP Alternatives

- Vector architectures
- SIMD extensions
- Graphics Processor Units (GPUs)

Supercomputer Applications

- Typical application areas
 - Military research (nuclear weapons, cryptography)
 - Scientific research
 - Weather forecasting
 - Oil exploration
 - Industrial design (car crash simulation)
 - Bioinformatics
 - Cryptography
- All involve huge computations on large data set
- Supercomputers: CDC6600, CDC7600, Cray-1, ...
- In 70s-80s, Supercomputer \equiv Vector Machine




Vector Supercomputers



[©Cray Research, 1976]

- Epitomized by Cray-1, 1976
- Scalar Unit
 - Load/Store Architecture
- Vector Extension
 - Vector Registers
 - Vector Instructions
- Implementation
 - Hardwired Control
 - Highly Pipelined Functional Units
 - Interleaved Memory System
 - No Data Caches
 - No Virtual Memory

Vector Programming Model



Vector Code Example

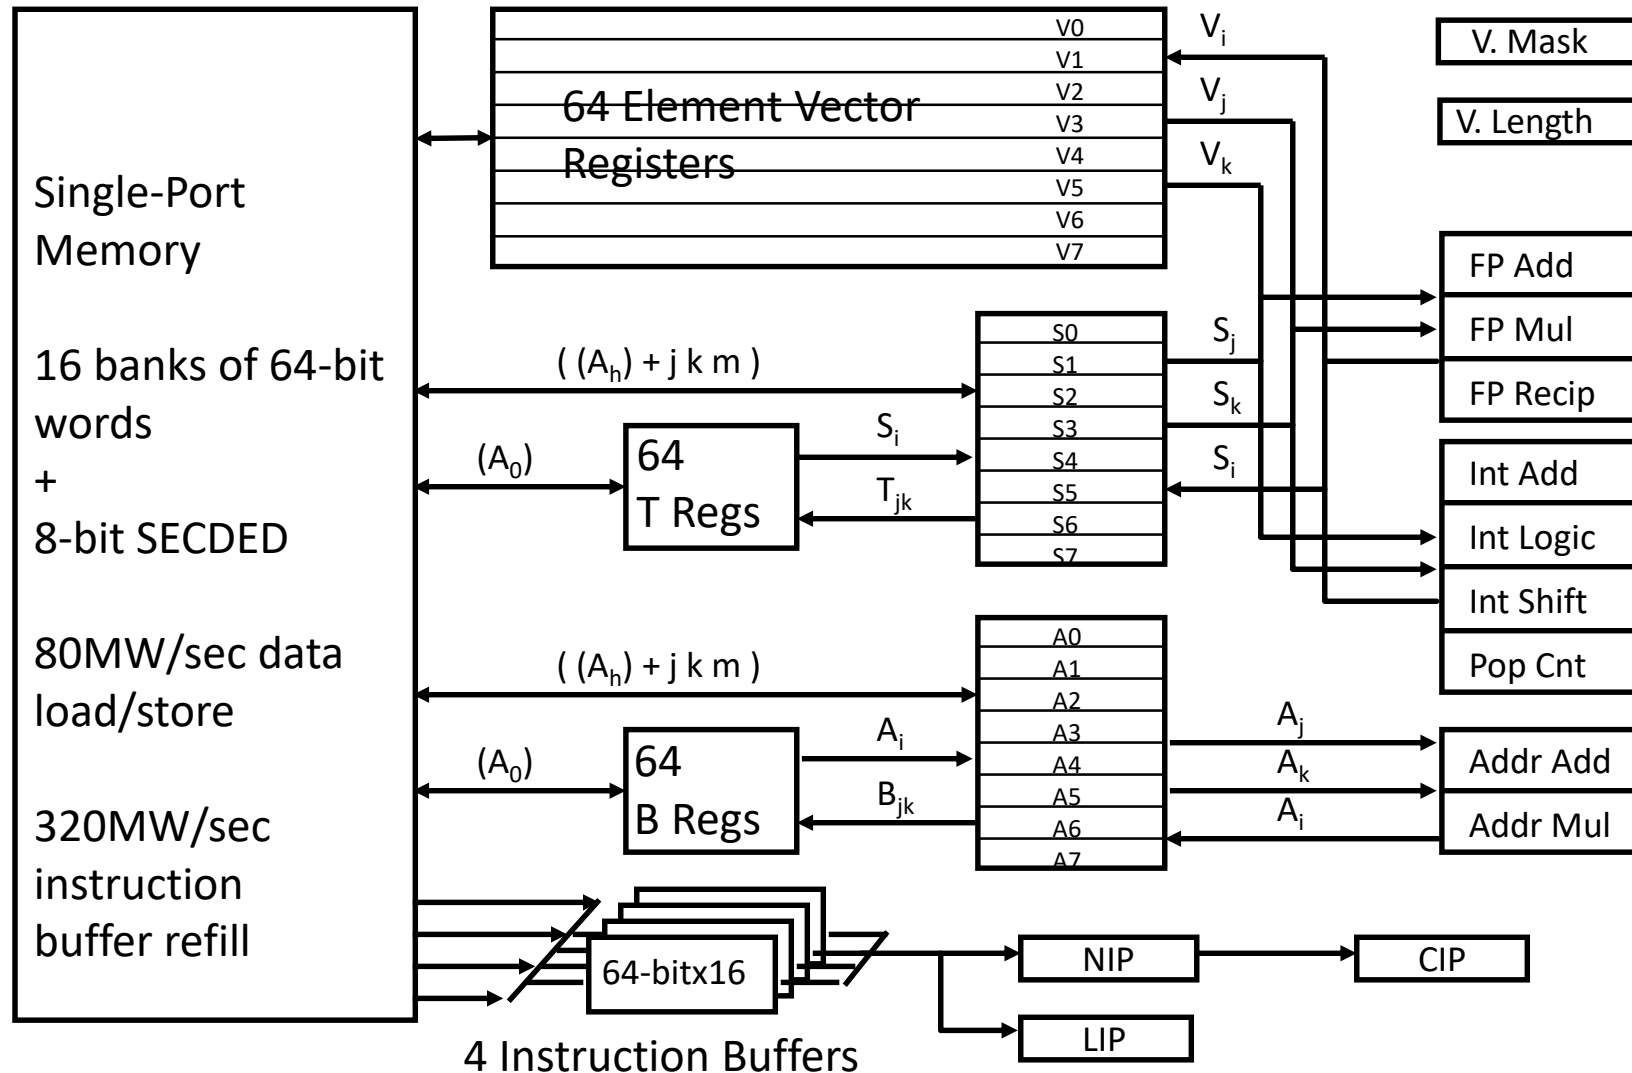
```
# C code
for (i=0; i<64; i++)
    C[i] = A[i] + B[i];
```

```
# Scalar Code
li x4, 64
loop:
    fld f1, 0(x1)
    fld f2, 0(x2)
    fadd.d f3, f1, f2
    fsd f3, 0(x3)
    addi x1, 8
    addi x2, 8
    addi x3, 8
    subi x4, 1
    bnez x4, loop
```

```
# Vector Code
li x4, 64
setv1 x4
vld v1, x1
vld v2, x2
vadd v3, v1, v2
vst v3, x3
```

区别在哪儿???

Cray-1 (1976)



- *memory bank cycle 50 ns*
- *processor cycle 12.5 ns (80MHz)*


Vector Instruction Set Advantages

- Compact
 - one short instruction encodes N operations
- Expressive, tells hardware that these N operations:
 - are independent
 - use the same functional unit
 - access disjoint registers
 - access registers in same pattern as previous instructions
 - access a contiguous block of memory (unit-stride load/store)
 - access memory in a known pattern (strided load/store)
- Scalable
 - can run same code on more parallel pipelines (lanes)

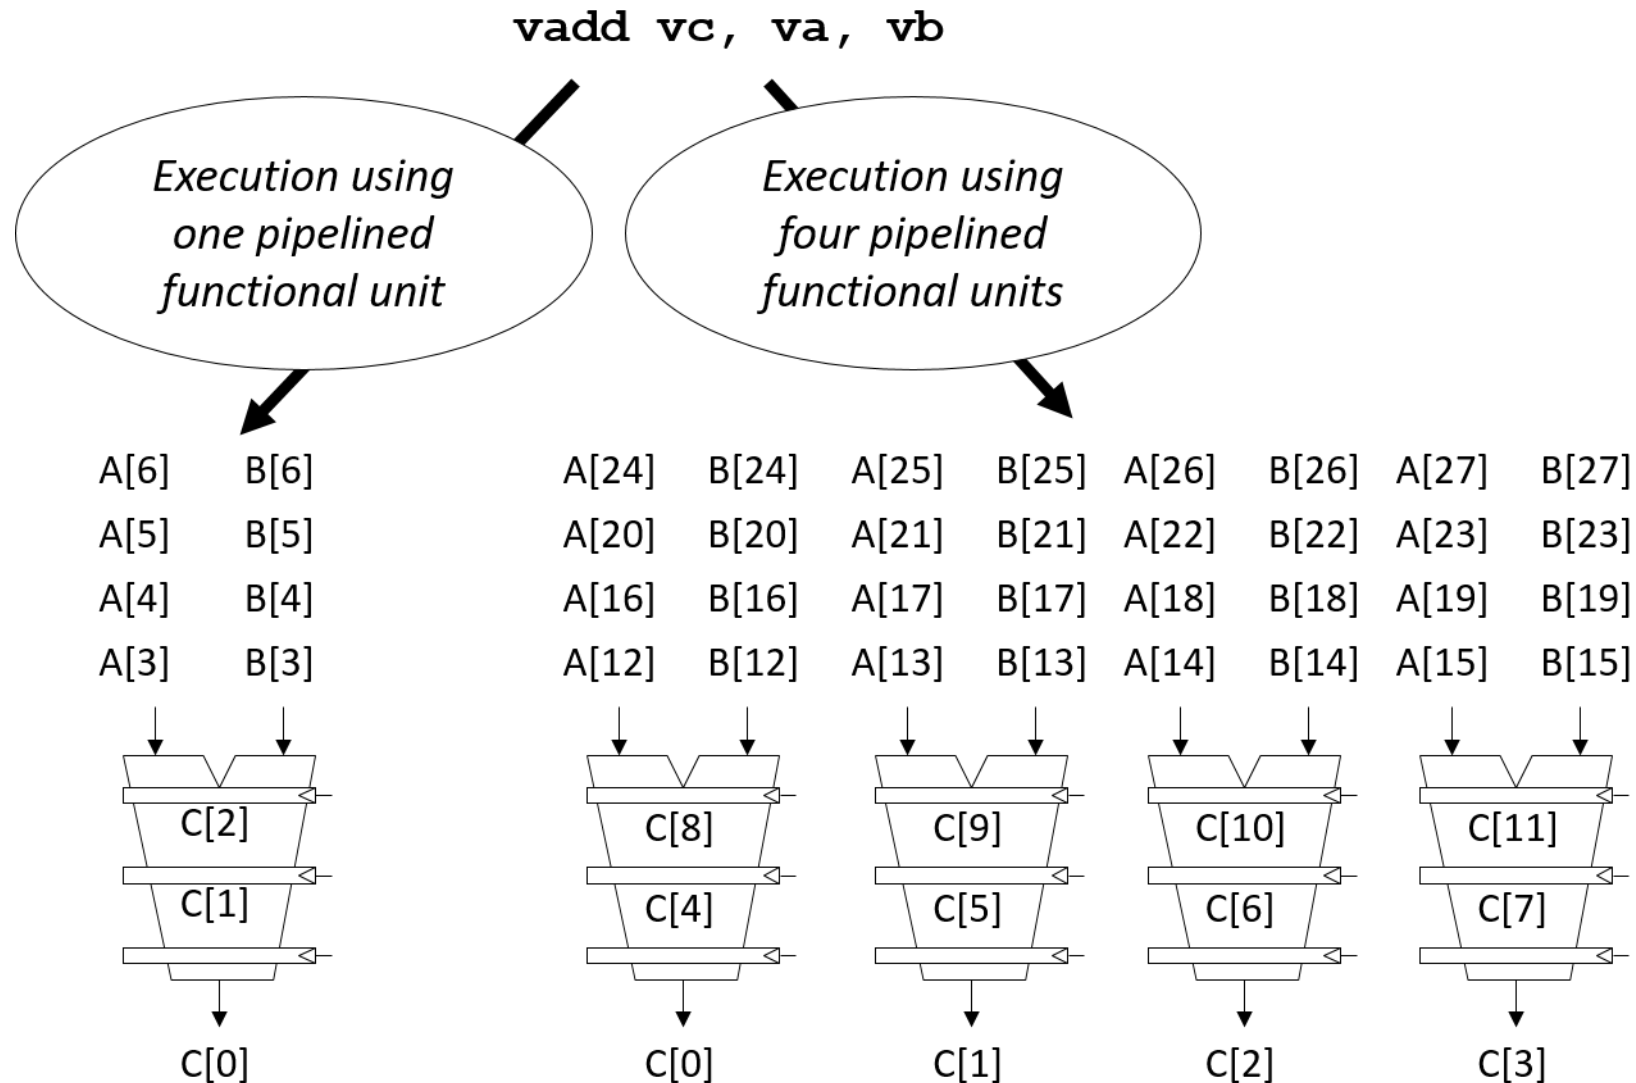
Vector Arithmetic Execution

- Use deep pipeline (\Rightarrow fast clock) to execute element operations
- Simplifies control of deep pipeline because elements in vector are independent (\Rightarrow no hazards!)

Six-stage multiply pipeline




Vector Instruction Execution




Interleaved Vector Memory System

- Bank busy time: Time before bank ready to accept next request
- Cray-1, 16 banks, 4 cycle bank busy time, 12 cycle latency

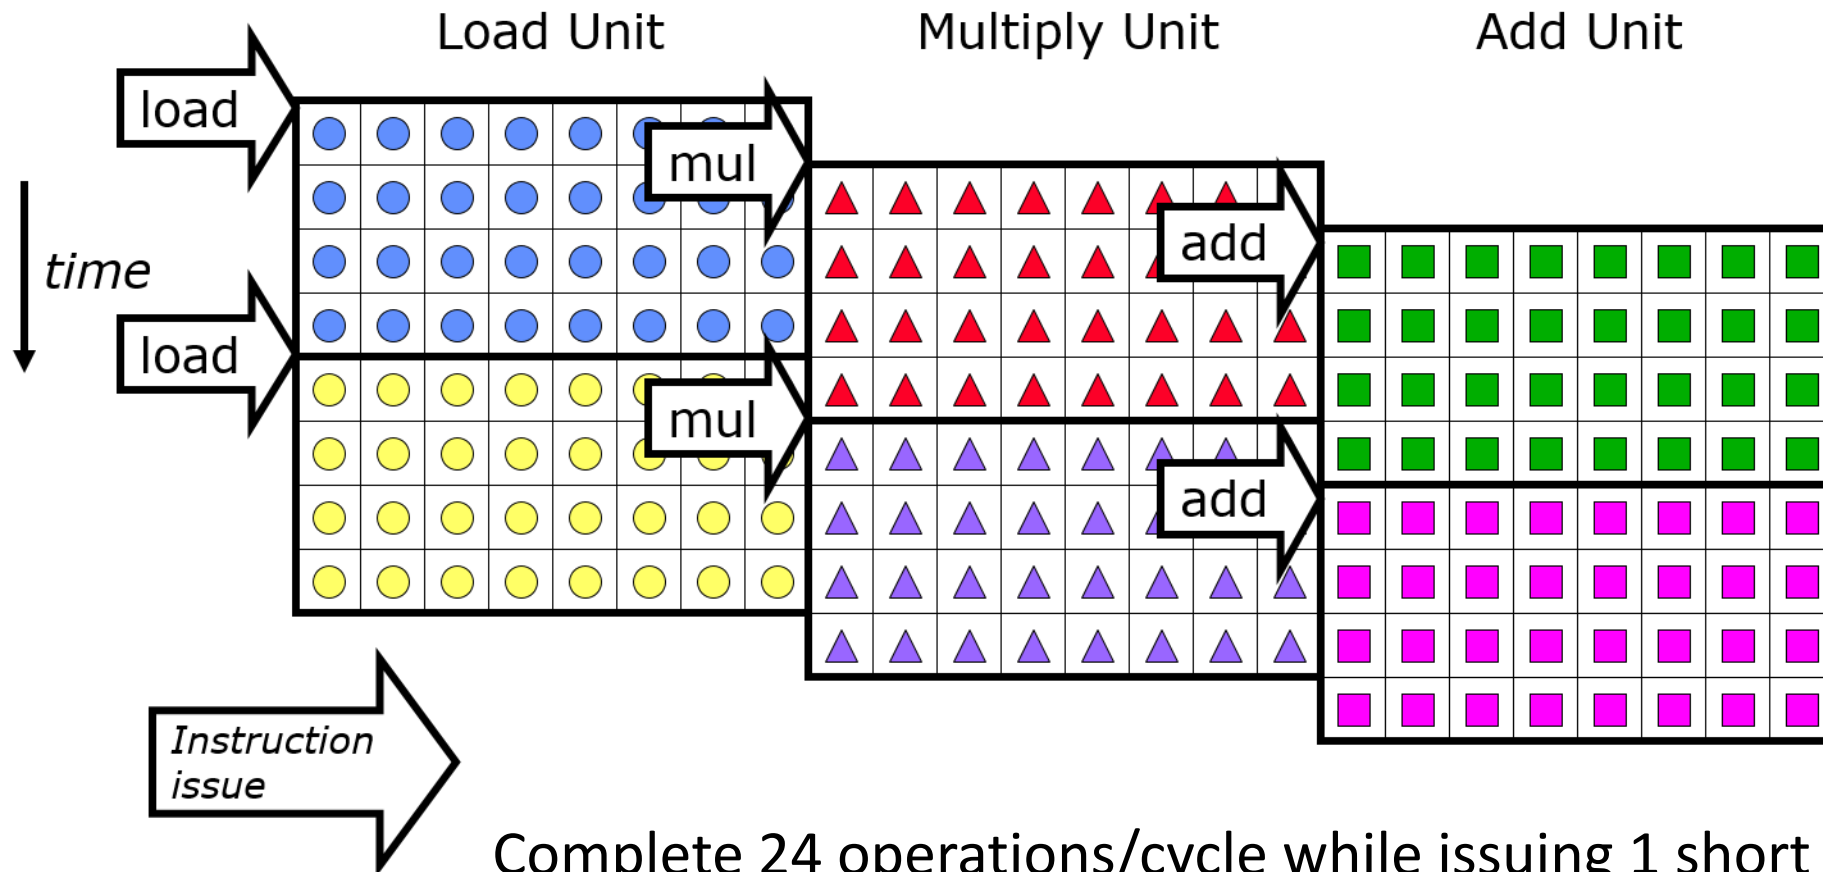


Vector Unit Structure



Vector Instruction Parallelism

- Can overlap execution of multiple vector instructions
 - example machine has 32 elements per vector register and 8 lanes




Complete 24 operations/cycle while issuing 1 short instruction/cycle

Vector Chaining

- Vector version of register bypassing
 - introduced with Cray-1

```
vld  v1  
vmul v3, v1, v2  
vadd v5, v3, v4
```




Vector Chaining Advantage

- Without chaining, must wait for last element of result to be written before starting dependent instruction




- With chaining, can start dependent instruction as soon as first result appears




Vector Startup


- Two components of vector startup penalty
 - functional unit latency (time through pipeline)
 - dead time or recovery time (time before another vector instruction can start down pipeline)



Dead Time and Short Vectors



No dead time →



T0, Eight lanes

No dead time

100% efficiency with 8-element vectors

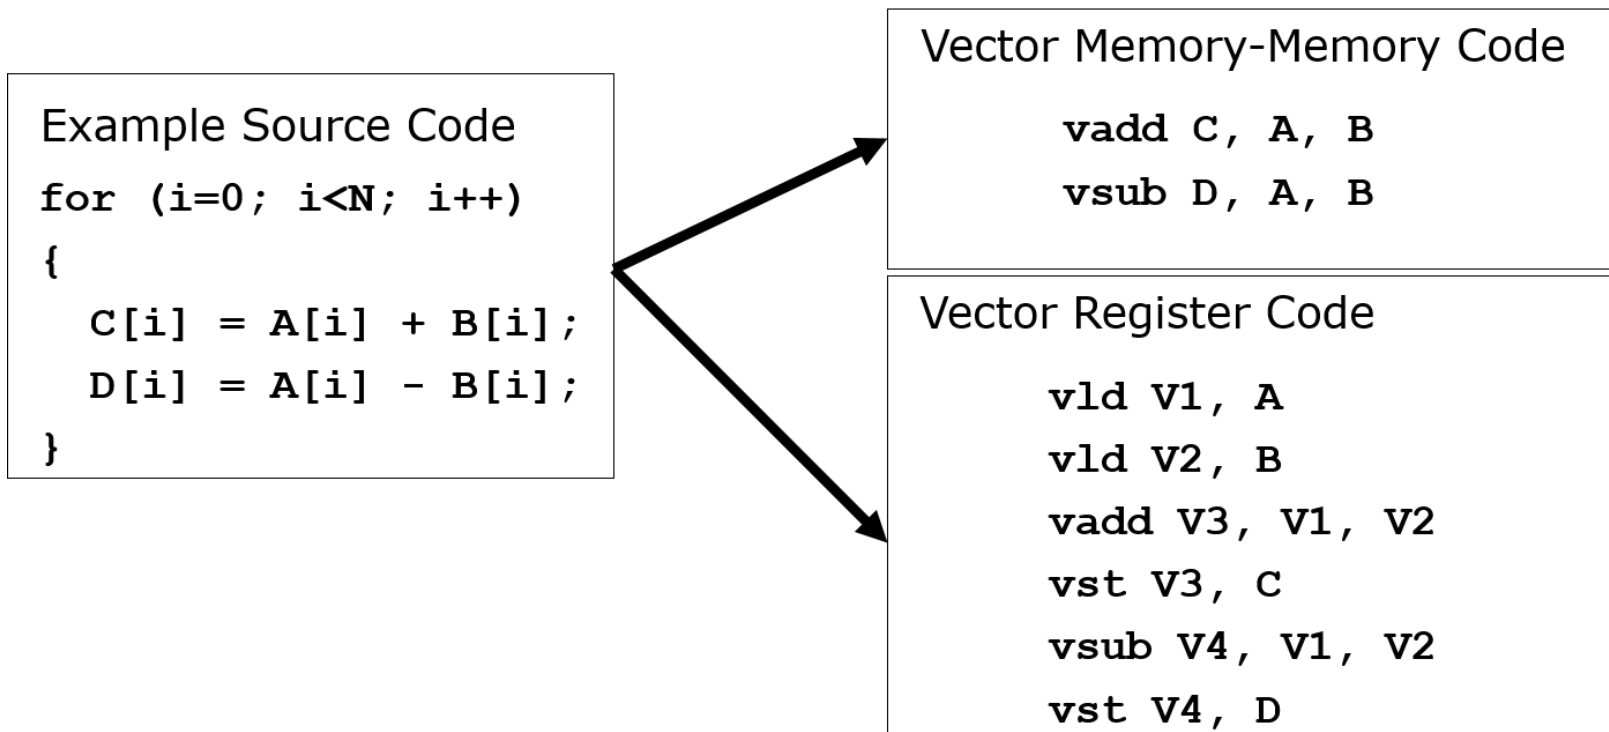
Cray C90, Two lanes

4-cycle dead time

Maximum efficiency 94% with 128-element vectors

Vector Memory-Memory versus Vector Register Machines

- Vector memory-memory instructions hold all vector operands in main memory
- The first vector machines, CDC Star-100 (1973) and TI ASC (1971), were memory-memory machines
- Cray-1 (1976) was first vector register machine



■ || Vector Memory-Memory versus Vector Register Machines

- Vector memory-memory architectures (VMMA) require greater main memory bandwidth, why?
 - All operands must be read in and out of memory
- VMMA make it difficult to overlap execution of multiple vector operations, why?
 - Must check dependencies on memory addresses
- VMMA incur greater startup latency
 - Scalar code was faster on CDC Star-100 for vectors < 100 elements
 - For Cray-1, vector/scalar breakeven point was around 2-4 elements
- Apart from CDC follow-ons (Cyber-205, ETA-10) all major vector machines since Cray-1 have had vector register architectures
- (we ignore vector memory-memory from now on)


Outline

- DLP
- SIMD
- GPU

Packed SIMD Extensions



- Short vectors added to existing ISAs for microprocessors
- Use existing 64-bit registers split into 2x32b or 4x16b or 8x8b
 - Lincoln Labs TX-2 from 1957 had 36b datapath split into 2x18b or 4x9b
 - Newer designs have wider registers
 - 128b for PowerPC AltiVec, Intel SSE2/3/4
 - 256b for Intel AVX
- Single instruction operates on all elements within register



SIMD in x86

- MMX (Multimedia Extension)
 - 1996, Pentium, 57 instructions
 - 64-bit MMX registers (borrowed from FP registers)
 - eight 8-bit / four 16-bit / two 32-bit operation
- SSE ~ SSE4 (Streaming SIMD Extension)
 - 1999~, Pentium III
 - 128-bit registers
 - FP: two 32-bit single precision, two 64-bit double precision
 - Int: four 32-bit, eight 16-bit, sixteen char
- AVX ~ AVX2 (Advanced Vector Extensions)
 - 2011~, Sandy Bridge
 - AVX: 256-bit registers, AVX2: 512-bit registers
 - Three-operand, relaxed alignment of memory operands

|| MMX

- After analyzing a lot of existing applications such as graphics, MPEG, music, speech recognition, game, image processing, they found that many multimedia algorithms execute the same instructions on many pieces of data in a large data set
- Typical elements are small, 8 bits for pixels, 16 bits for audio, 32 bits for graphics and general computing
- New data type: 64-bit packed data type. Why 64 bits?
 - Good enough
 - Practical

MMX Data Types

Packed Byte: 8 bytes packed into 64 bits



Packed Word: 4 words packed into 64 bits




Packed Doubleword: 2 doublewords packed into 64 bits



Packed Quadword: One 64-bit quantity



MMX Registers



8 MMX Registers MM0 ~ MM7

No new mode or state was created. Hence, for context switching, no extra state needs to be saved.

It means MMX and FPU cannot be used at the same time. Big overhead to switch

MMX Instructions

Category		Wraparound	Signed Saturation	Unsigned Saturation
Arithmetic	Addition	PADDB, PADDW, PADDD	PADDSB, PADDSW	PADDUSB, PADDUSW
	Subtraction	PSUBB, PSUBW, PSUBD	PSUBSB, PSUBSW	PSUBUSB, PSUBUSW
	Multiplication Multiply and Add	PMULL, PMULH PMADD		
Comparison	Compare for Equal	PCMPEQB, PCMPEQW, PCMPEQD		
	Compare for Greater Than	PCMPGTPB, PCMPGTPW, PCMPGTPD		
Conversion	Pack		PACKSSWB, PACKSSDW	PACKUSWB
Unpack	Unpack High	PUNPCKHBW, PUNPCKHWD, PUNPCKHDQ		
	Unpack Low	PUNPCKLBW, PUNPCKLWD, PUNPCKLDQ		


MMX Instructions

		Packed	Full Quadword
Logical	And And Not Or Exclusive OR		PAND PANDN POR PXOR
Shift	Shift Left Logical Shift Right Logical Shift Right Arithmetic	PSLLW, PSLLD PSRLW, PSRLD PSRAW, PSRAD	PSLLQ PSRLQ
Data Transfer	Register to Register Load from Memory Store to Memory	Doubleword Transfers	Quadword Transfers
		MOVD MOVD MOVD	MOVQ MOVQ MOVQ
Empty MMX State		EMMS	


Call it before you switch to FPU from MMX;
Expensive operation

■ ■ ■ Saturation arithmetic

- Useful in Graphics Applications



wrap-around



saturating


Example: add a constant to a vector

```
char d[]={5, 5, 5, 5, 5, 5, 5, 5};  
char clr[]={65,66,68,...,87,88}; // 24 bytes  
__asm{  
    movq mm1, d  
    mov cx, 3  
    mov esi, 0  
L1: movq mm0, clr[esi]  
    paddb mm0, mm1  
    movq clr[esi], mm0  
    add esi, 8  
    loop L1  
    emms  
}
```



Application: frame difference




|A-B|



Application: frame difference



(A-B) or (B-A)



Application: frame difference

```
MOVQ      mm1, A //move 8 pixels of image A
MOVQ      mm2, B //move 8 pixels of image B
MOVQ      mm3, mm1 // mm3=A
PSUBSB    mm1, mm2 // mm1=A-B
PSUBSB    mm2, mm3 // mm2=B-A
POR       mm1, mm2 // mm1=|A-B|
```

Example: image fade-in-fade-out



A




B

$$A * \alpha + B * (1 - \alpha) = B + \alpha(A - B)$$

|| $\alpha=1$



|| $\alpha=0.75$



|| $\alpha=0.5$



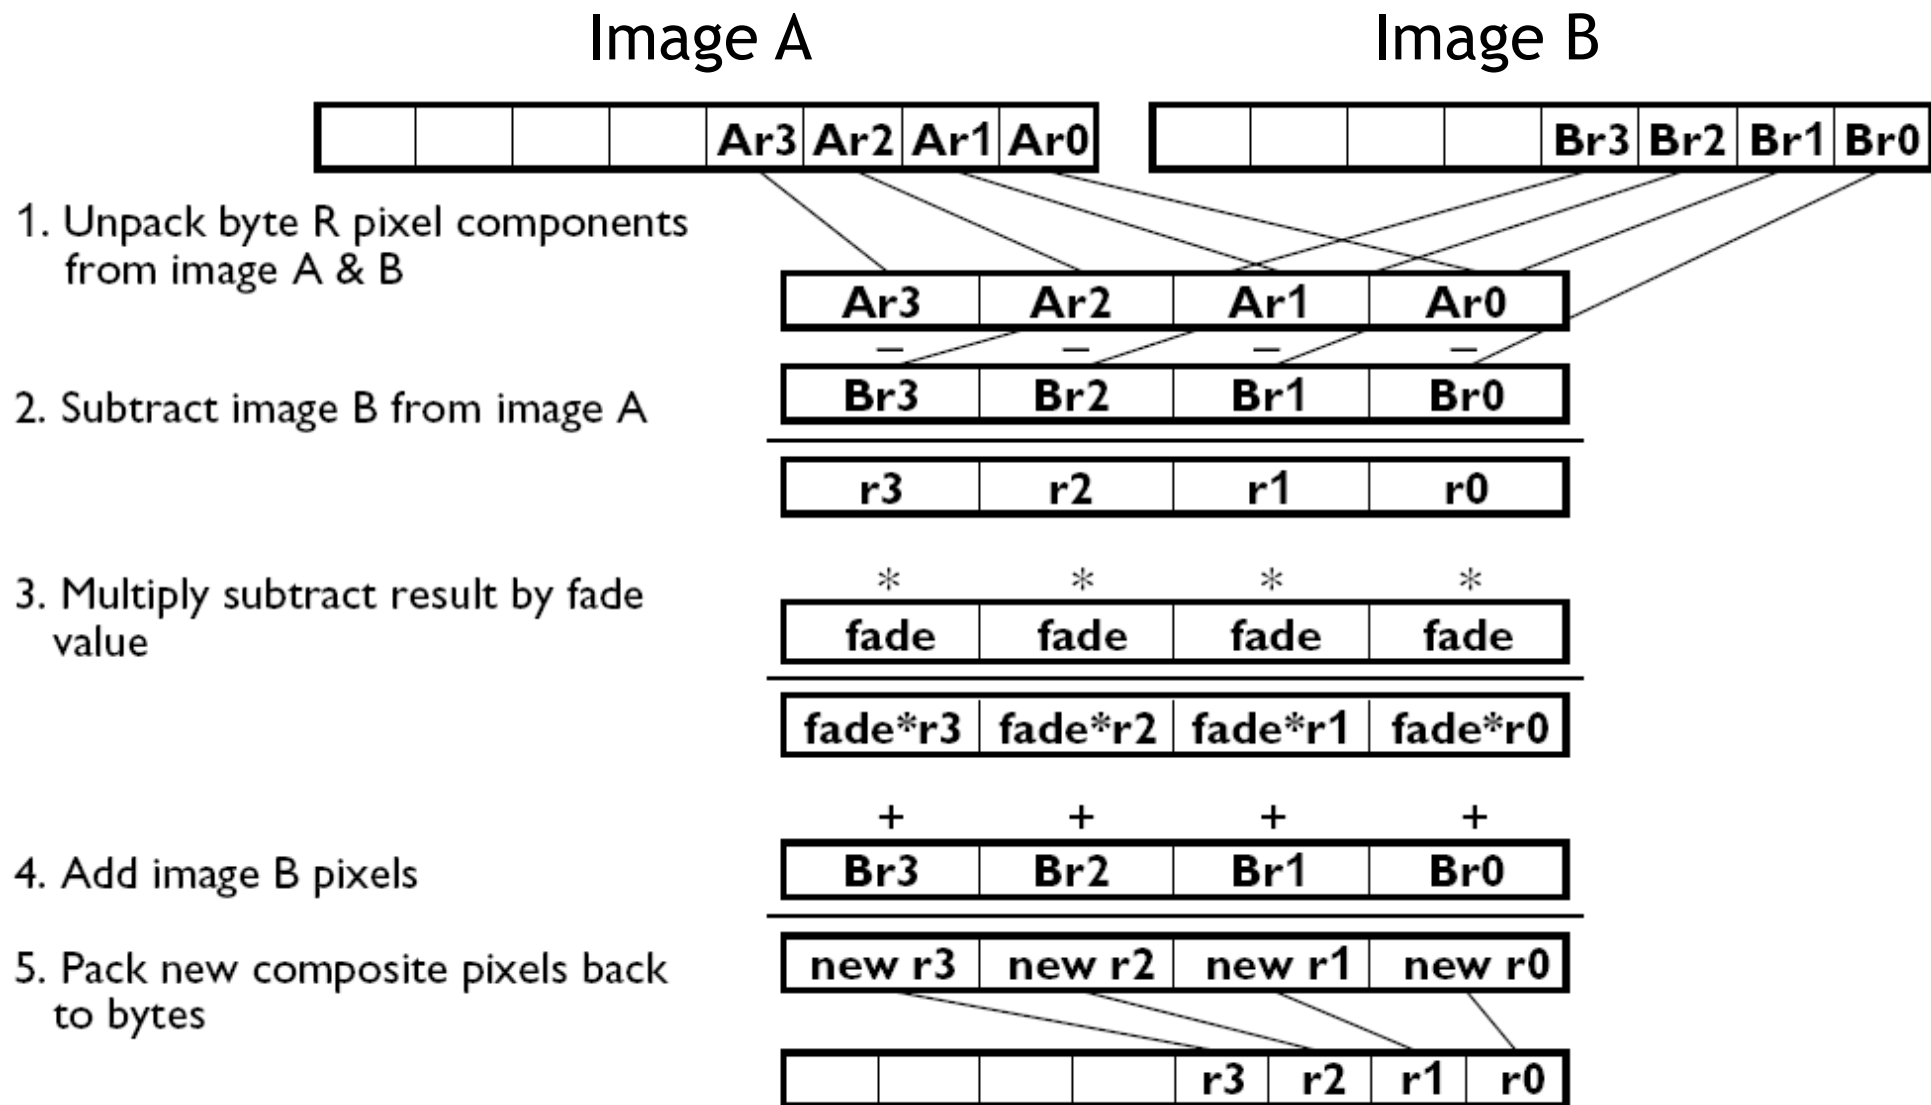
|| $\alpha=0.25$



||| $\alpha=0$



Example: image fade-in-fade-out



Example: image fade-in-fade-out

```

MOVQ      mm0, alpha//4 16-b zero-padding  $\alpha$ 
MOVD      mm1, A //move 4 pixels of image A
MOVD      mm2, B //move 4 pixels of image B
PXOR      mm3, mm3 //clear mm3 to all zeroes
//unpack 4 pixels to 4 words
PUNPCKLBW mm1, mm3 // Because B-A could be
PUNPCKLBW mm2, mm3 // negative, need 16 bits
PSUBW     mm1, mm2 //(B-A)
PMULHW    mm1, mm0 //(B-A)*fade/256
PADDW     mm1, mm2 //(B-A)*fade + B
//pack four words back to four bytes
PACKUSWB  mm1, mm3
    
```

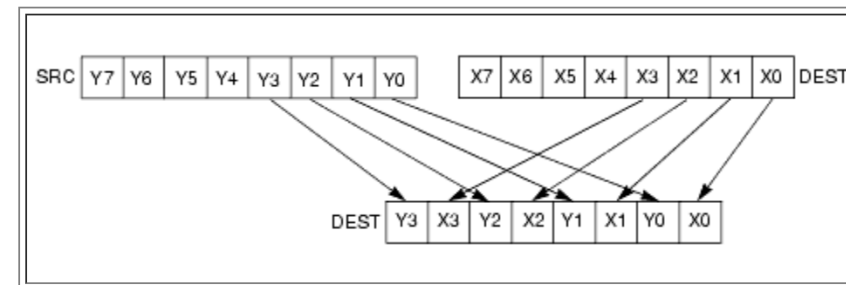





图 3-15. PUNPCKLBW 指令操作

PACKUSWB mm, mm/m64




Application: matrix transpose

Elements: 16-bit integers




Little endian: LSB => Lower address




Application: matrix transpose

Phase 1



Phase 2




Note: Repeat for the other rows to generate $([d_3, c_3, b_3, a_3]$ and $[d_2, c_2, b_2, a_2])$.

Application: matrix transpose

MMX code sequence operation:


movq	mm1, row1	; load pixels from first row of matrix
movq	mm2, row2	; load pixels from second row of matrix
movq	mm3, row3	; load pixels from third row of matrix
movq	mm4, row4	; load pixels from fourth row of matrix
punpcklwd	mm1, mm2	; unpack low order words of rows 1 & 2, mm1 = [b1, a1, b0, a0]
punpcklwd	mm3, mm4	; unpack low order words of rows 3 & 4, mm3 = [d1, c1, d0, c0]
movq	mm5, mm1	; copy mm1 to mm5
punpckldq	mm1, mm3	; unpack low order doublewords -> mm2 = [d0, c0, b0, a0]
punpckhdq	mm5, mm3	; unpack high order doublewords -> mm5 = [d1, c1, b1, a1]



Application: matrix transpose

MMX code sequence operation:


movq	mm1, row1	; load pixels from first row of matrix
movq	mm2, row2	; load pixels from second row of matrix
movq	mm3, row3	; load pixels from third row of matrix
movq	mm4, row4	; load pixels from fourth row of matrix
punpcklwd	mm1, mm2	; unpack low order words of rows 1 & 2, mm1 = [b1, a1, b0, a0]
punpcklwd	mm3, mm4	; unpack low order words of rows 3 & 4, mm3 = [d1, c1, d0, c0]
movq	mm5, mm1	; copy mm1 to mm5
punpckldq	mm1, mm3	; unpack low order doublewords -> mm2 = [d0, c0, b0, a0]
punpckhdq	mm5, mm3	; unpack high order doublewords -> mm5 = [d1, c1, b1, a1]




Application: matrix transpose

MMX code sequence operation:


```
movq      mm1, row1      ; load pixels from first row of matrix
movq      mm2, row2      ; load pixels from second row of matrix
movq      mm3, row3      ; load pixels from third row of matrix
movq      mm4, row4      ; load pixels from fourth row of matrix
punpcklwd mm1, mm2        ; unpack low order words of rows 1 & 2, mm1 = [b1, a1, b0, a0]
punpcklwd mm3, mm4        ; unpack low order words of rows 3 & 4, mm3 = [d1, c1, d0, c0]
movq      mm5, mm1        ; copy mm1 to mm5
punpckldq mm1, mm3        ; unpack low order doublewords -> mm2 = [d0, c0, b0, a0]
punpckhdq mm5, mm3        ; unpack high order doublewords -> mm5 = [d1, c1, b1, a1]
```




mm3



mm5



mm3

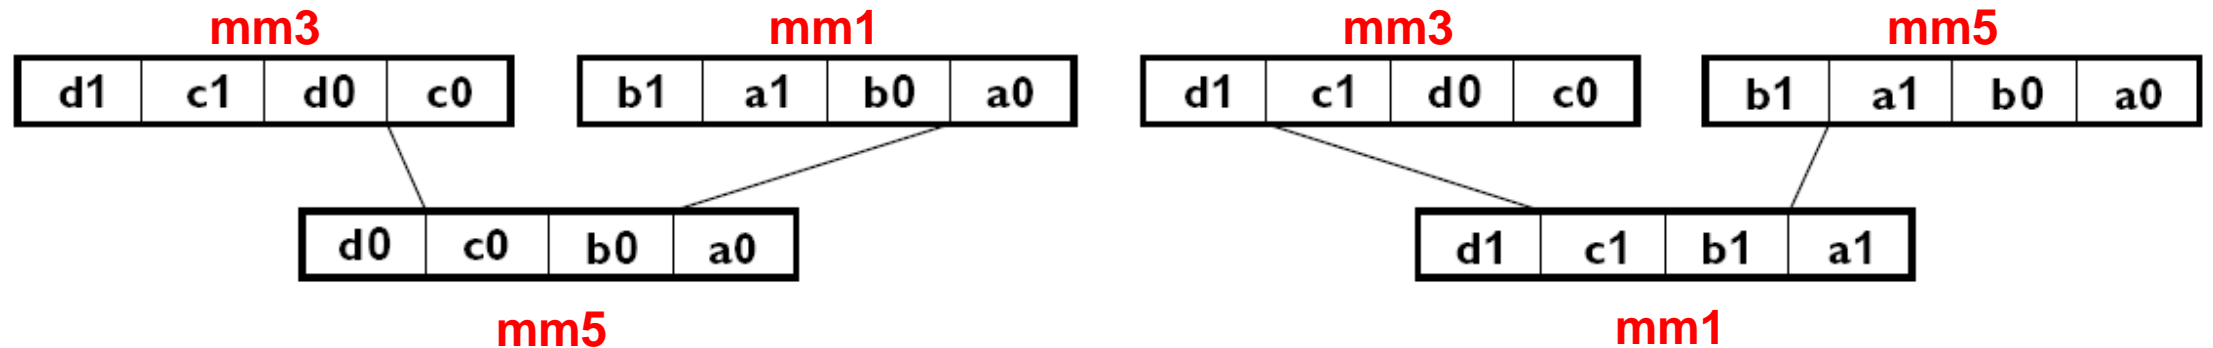


mm1

Application: matrix transpose


MMX code sequence operation:

```
movq      mm1, row1      ; load pixels from first row of matrix
movq      mm2, row2      ; load pixels from second row of matrix
movq      mm3, row3      ; load pixels from third row of matrix
movq      mm4, row4      ; load pixels from fourth row of matrix
punpcklwd mm1, mm2        ; unpack low order words of rows 1 & 2, mm1 = [b1, a1, b0, a0]
punpcklwd mm3, mm4        ; unpack low order words of rows 3 & 4, mm3 = [d1, c1, d0, c0]
movq      mm5, mm1        ; copy mm1 to mm5
punpckldq mm1, mm3        ; unpack low order doublewords -> mm2 = [d0, c0, b0, a0]
punpckhdq mm5, mm3        ; unpack high order doublewords -> mm5 = [d1, c1, b1, a1]
```




Performance Boost

- Benchmark kernels: FFT, FIR, vector dot-product, IDCT, motion compensation.
- >80% performance gain
- Lower the cost of multimedia programs by removing the need of specialized DSP chips



DLP important for conventional CPUs



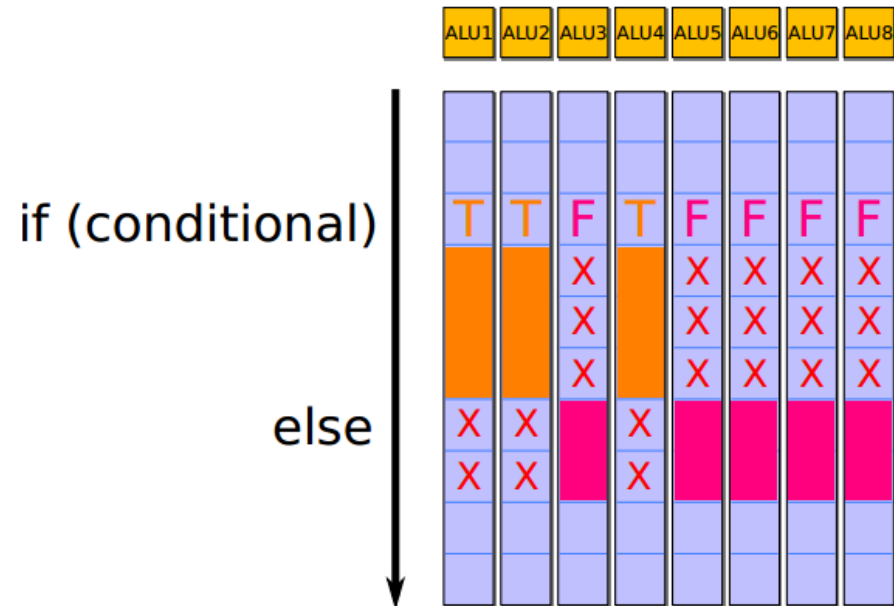
- Prediction for x86 processors, from Hennessy & Patterson, 5th edition
 - *Note: Educated guess, not Intel product plans!*
- TLP: 2+ cores / 2 years
- DLP: 2x width / 4 years
- DLP will account for more mainstream parallelism growth than TLP in next decade.
 - SIMD –single-instruction multiple-data (DLP)
 - MIMD- multiple-instruction multiple-data (TLP)

SIMD Drawbacks

- Restricted degree of parallelism: 64-bit ~ 512-bit
 - Continuous extension to ISA?
- Restricted conditions
 - Effective on all elements, no flexibility on conditional cases
- Problem: Want to vectorize loops with conditional code:

```
for (i=0; i<N; i++)  
    if (A[i]>0) then  
        A[i] = B[i];
```

SIMD \neq DLP !



■ ■ ■ Packed SIMD versus Vectors

- Limited instruction set
 - no vector length control
 - no strided load/store or scatter/gather
 - unit-stride loads must be aligned to 64/128-bit boundary
- Limited vector register length:
 - requires superscalar dispatch to keep multiply/add/load units busy
 - loop unrolling to hide latencies increases register pressure
- Trend towards fuller vector support in microprocessors
 - Better support for misaligned memory accesses
 - Support of double-precision (64-bit floating-point)
 - New Intel AVX spec (announced April 2008), 256b vector registers (expandable up to 1024b) , adding scatter/gather
 - New ARM SVE/MVE vector ISA closer to traditional vector designs

Outline

- DLP
- SIMD
- GPU

■ ■ ■ Graphics Processing Units (GPUs)

- Original GPUs were dedicated fixed-function devices for generating 3D graphics (mid-late 1990s) including high-performance floating-point units
 - Provide workstation-like graphics for PCs
 - User could configure graphics pipeline, but not really program it
- Over time, more programmability added (2001-2005)
 - E.g., New language Cg for writing small programs run on each vertex or each pixel, also Windows DirectX variants
 - Massively parallel (millions of vertices or pixels per frame) but very constrained programming model
- Some users noticed they could do general-purpose computation by mapping input and output data to images, and computation to vertex and pixel shading computations
 - Incredibly difficult programming model as had to use graphics pipeline model for general computation

■ ■ ■ General-Purpose GPUs (GP-GPUs)

- In 2006, Nvidia introduced GeForce 8800 GPU, which supported a new programming language CUDA (in 2007)
 - “Compute Unified Device Architecture”
 - Subsequently, broader industry pushing for OpenCL, a vendor-neutral version of same ideas.
- Idea: Take advantage of GPU computational performance and memory bandwidth to accelerate some kernels for general-purpose computing
- Attached processor model: Host CPU issues data-parallel kernels to GP-GPU for execution

Early 90s – Pre GPU



Wolfenstein 3D, 1992





Doom I, 1993

- Interactive software rendering (no GPUs yet)
- NOTE: SGI was building interactive rendering supercomputers, but this was beginning of interactive 3D graphics on PC


Graphics Workloads

- Triangles/vertices and pixels/fragments




CPU and GPU Trends

Theoretical GFLOP/s



CPU and GPU Trends





NIPS 2017 NVIDIA 发布最新GPU (Titan V)

- Volta 架构 GV100, 拥有 12GB 显存
- 峰值浮点性能为 110 Tflops, 是Titan Xp (12T) 的 9 倍
- 在 fp32 下是 Titan Xp 的 1.5 倍, 在 fp16 下是 Titan Xp 的 3 倍
- 售价 \$2999 (大概2w RMB)
- 采用台积电 (TSMC) 12nm FFN 制程
- 5120 个 CUDA 核心
- 晶体管数量: 211 亿
- 面积: 815mm^2 (约等于一块apple watch的面积)

Important of Machine Learning for GPUs

NVIDIA Corporation

NVDA

\$501.41 ↑30,473.78% +499.77 MAX


After Hours: **\$501.50** (↑0.018%) +\$0.090

Closed: Mar 25, 5:46:41 PM UTC-4 · USD · NASDAQ · Disclaimer

1D 5D 1M 6M YTD 1Y 5Y **MAX**





NVIDIA stock price 40x in 9 years (since deep learning became important)



欢迎提问

Vector vs. SIMD



■ ■ ■ Vector vs. GPU

- GPU
 - Discrete accelerator running in own memory space
 - More recently support sharing of physical, or pinned virtual, but still in different memory hierarchy and no page faults
 - Terrible at scalar code
 - Only effective on very large data-parallel tasks (>10,000s)
 - ISA/microarchitecture evolved from graphics shader needs, not general compute
 - Only seems efficient compared to out-of-order scalar core
- Traditional vector
 - Coprocessor running in same memory hierarchy as scalar
 - Tightly coupled to scalar core
 - Effective with loop counts of 2 or more
 - ISA evolved from general computing needs
 - Very efficient